

Computer Vision and Machine Learning

(Neural Network-1)

Bhabatosh Chanda
bchanda57@gmail.com

Regression

- Regression is a technique to establish relation between independent variables or primary observations (features) and dependent variables.
- Depending on type of relation between dependent variable (*decision or prediction*) and independent variable(s), we may classify the regression as
 - *Linear regression*
 - *Logistic regression*

Intro 2 ML

2

Machine learning tasks T

- **Classification:** To decide which of the *k classes* the given input belongs to. Learning system tries to develop a mapping (function)

$$f: R^n \rightarrow \{1, 2, \dots, k\}$$

- **Prediction:** To predict a numerical value for the given input. So the task is similar to classification except the representation of output. Thus the mapping (function) is

$$f: R^n \rightarrow R$$

Intro 2 ML

3

Idea of machine learning

- A system (here, machine or computer) is said to have learned
 - to do some *task T*
 - from a set of *examples E*
 - in terms of a *performance measure P* ,


if its performance improves

- as measured by the same *P*
- to carry out the same task *T*
- by dealing with the example set *E* .

Intro 2 ML

4

Machine learning vs. Deep learning



Intro 2 ML

5

Linear regression

- Task is to build a system to predict a scalar value $y \in R$ as output from the given input $x \in R^n$.
- Suppose \hat{y} is the value predicted by the system, i.e.,

$$\hat{y} = \mathbf{w}^T \mathbf{x}$$

where $\mathbf{w} \in R^n$ is parameter vector that controls behaviour of system.

- Assume $\mathbf{x} = (x_1, x_2, \dots, x_n)$, similarly $\mathbf{w} = (w_1, w_2, \dots, w_n)$

Intro 2 ML

6

Linear regression (contd.)

- A more general relation between $x \in R^n$ and $y \in R$ may be expressed as

$$\hat{y} = \mathbf{w}^T \mathbf{x} + b$$


- If we append a '1' to \mathbf{x} and including ' b ' as a weight
 - Relation between y and \mathbf{x} becomes affine, but
 - Relation between y and \mathbf{w} remains linear.
- Consider $\mathbf{x} = (x_0, x_1, x_2, \dots, x_n)$ and $\mathbf{w} = (w_0, w_1, w_2, \dots, w_n)$ where $x_0 = 1$ and $w_0 = b$.

Intro 2 ML

7

Linear regression (contd.)

Example:



4/10/2024

8

Linear regression (contd.)

- $x \rightarrow$ independent variable (e.g., age of a deer, time in quarter, etc.)
- $y \rightarrow$ dependent variable (resp., weight of a deer, pairs of shoes sold)
- Let us consider relation between x and y may be modeled as a straight line:

$$y = wx + b$$


- Exploiting linear regression technique, we estimate

$$w = \frac{\sum(x - \bar{x})(y - \bar{y})}{\sum(x - \bar{x})^2} \quad b = \bar{y} - w\bar{x}$$

4/10/2024

9


Linear regression (contd.)



4/10/2024

10

Prediction



4/10/2024

11

Prediction: multiple input

- So far we have discussed the cases where input is a single variable.

$$y = f(x)$$

- No. of input variables (independent variables) may be more than 1.

$$y = f(x_1, x_2, x_3, \dots, x_n)$$


- A contrived example may have following input variables:

Variable	1	2	3	4	5	6	7	8	9	10
x_1	37	42	38	34	41	42	36	40	39	43
x_2	95	93	97	96	98	98	94	97	99	95
y	0	0	0	0	1	1	0	1	1	1

4/10/2024

12

Binary classification: Pomfret and Magur



4/10/2024

13

Two class problem: Pomfret and Magur



4/10/2024

14


Features: Pomfret and Magur



4/10/2024

15

Two-class problem: Feature space



4/10/2024

16

Boundary function

- Find coefficients a , b and c of equation of a straight line

$$ax + by + c = 0$$

such that for all observation a feature pair (x, y) :

$$ax + by + c > 0 \quad \text{if } (x, y) \text{ belongs to } C_1$$


$$ax + by + c < 0 \quad \text{if } (x, y) \text{ belongs to } C_2$$

- If the desired condition is not satisfied for any feature-label pair we call a classification error has occurred.
- In general, decision boundary must be estimated to minimize this error.

4/10/2024

17


Two-class problem



4/10/2024

18

Two-class problem



4/10/2024

19

Generalization

- The ability to perform well on previously unseen data is called **generalization**.
- The target of machine learning to keep **generalization error** or **test error** as low as possible.
 - Note that system is built by minimizing the train error.
 - Is there any relation between training error and test error?

Intro 2 ML

20

Generalization (contd.)

- Training and test data are accumulated by same data generating process.
 - Each example in training and test datasets are **independent** to each other.
 - The training and test datasets are **identically distributed**.
- The *i.i.d.* assumption allows us to study the relationship between the training error and the test error.
 - Expected training error and the expected test error of a model are equal.

Intro 2 ML

21


Overfitting and underfitting

- **Two criteria** that determines how well a machine learning algorithm performs are its ability to
 1. **make the training error small, and**
 2. **Make the gap between the training error and the test error small.**
- These correspond to two problems: **overfitting** and **underfitting**.
 - If the training error is not small \rightarrow underfitting
 - If gap between training and test errors is not small \rightarrow overfitting.

Intro 2 ML

22


Overfitting and underfitting (contd.)



Intro 2 ML

23

Overfitting and underfitting (contd.)



Intro 2 ML

24

How to set the boundary function

- Based on the training data set.
 - All at a time.
 - Linear discriminant analysis
 - One at a time.
 - Perceptron network, neural network

Intro 2 ML

25


How to set the boundary function

- Based on the training data set.
 - All at a time.
 - Linear discriminant analysis
 - One at a time.
 - Perceptron network, neural network

Intro 2 ML

26


Forming the decision boundary



Intro 2 ML

27



Forming the decision boundary



Intro 2 ML

28



Forming the decision boundary



Intro 2 ML

29



Forming the decision boundary



Intro 2 ML

30



Forming the decision boundary



Intro 2 ML

31


Forming the decision boundary



Intro 2 ML

32


Forming the decision boundary



Intro 2 ML

33


Forming the decision boundary



Intro 2 ML

34


Forming the decision boundary



Intro 2 ML

35


Forming the decision boundary



Intro 2 ML

36


Forming the decision boundary



Intro 2 ML

37


Forming the decision boundary



Intro 2 ML

38


Forming the decision boundary



Intro 2 ML

39


Forming the decision boundary



Intro 2 ML

40


Forming the decision boundary



Intro 2 ML

41


Forming the decision boundary



Intro 2 ML

42


Forming the decision boundary



Intro 2 ML

43

Forming the decision boundary



Intro 2 ML

44

Boundary function

- Find coefficients a , b and c of equation of a straight line

$$ax + by + c = 0$$

such that for all observation a feature pair (x, y) :

$$ax + by + c > 0 \quad \text{if } (x, y) \text{ belongs to } C_1$$



$$ax + by + c < 0 \quad \text{if } (x, y) \text{ belongs to } C_2$$

- If the desired condition is not satisfied for any feature-label pair we call a classification error has occurred.
- In general, decision boundary must be estimated to minimize this error.

4/10/2024

45

Linear classifier and neuron



4/10/2024

46



What are Artificial Neural Networks?

- Mimics the function of the brain and nervous system
- Highly parallel
 - Process information much more like the brain than a serial computer
- Learning
- Very simple principles
- Very complex behaviours

4/10/2024

47

Neuron versus Node




4/10/2024

48

Function of a node

- At node
Output $O = f(\sum w_i x_i)$
where $f(\cdot)$ is a squashing function.
- Squashing function limits node output.




4/10/2024

49

Perceptron

- Linear threshold unit (LTU)




50

Perceptron network

- Synonym for single layer, feed-forward network capable of learning.
- Output $O = f(\sum_j W_j I_j + b_j)$

where 'b' is bias, which however, may be included as additional weight.




4/10/2024

51

Feed-forward nets

- Information flow is unidirectional
 - Data is presented to *Input layer*
 - Passed on to *Hidden Layer*
 - Passed on to *Output layer*
- Information is distributed
- Information processing is parallel
- True while testing new data



4/10/2024

52

Standard activation functions

- The hard-limiting threshold function
 - Corresponds to the biological paradigm
 - either fires or not (**Perceptron**)
- Sigmoid functions ('S'-shaped curves)
 - The hyperbolic tangent (symmetrical)
 - Both functions have a simple differential
 - Only the shape is important (**Neuron**)

$$\phi(x) = \frac{1}{1 + e^{-ax}}$$

4/10/2024


53

Example: node function

- Feeding data through the net:

$$(1 \times 0.25) + (0.5 \times (-1.5)) = 0.25 + (-0.75) = -0.5$$

Squashing: $\frac{1}{1 + e^{0.5}} = 0.3775$



4/10/2024

54

Loss function or Error or Cost function

- Training sample is composed of
 - Input data (feature vector) and
 - Actual class label (also known as groundtruth)
- Given the input, feed forward network predicts class label
 - based on current parameters
 - Loss or error or cost is measured as total deviation from groundtruth

Cost or Loss or Error: $E(\mathbf{w}) = \sum (\text{Predicted label} - \text{Actual label})^2$
 where \mathbf{w} is parameter vector.


4/10/2024

55

Training the network

- Means setting correct weights (including bias) or parameters of the network.
- Backpropagation
 - Requires training set (input / output pairs)
 - Starts with small random weights
 - Compute error between predicted label and actual label (groundtruth)
 - Error is used to adjust weights (supervised learning)

→ Gradient descent on error landscape



4/10/2024

56

Machine learning network

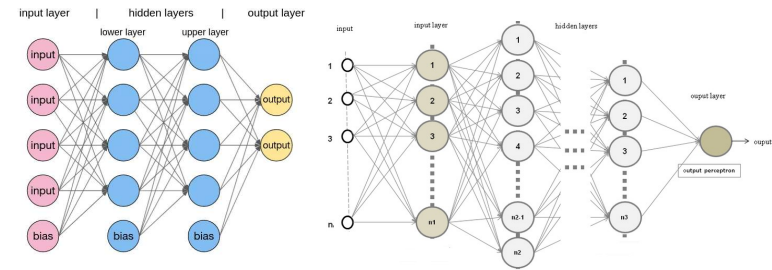
Machine learning models for classification have followings are common:

- **Input layer:** quantitative representation of object features
- **Hidden layer(s):** apply transformations with nonlinearity
- **Output layer:** Result for classification, regression etc.
- The models are trained through **supervised learning**.
 - Training data are explicitly labelled (known output).
 - Weights are updated to minimize error between prediction and the groundtruth.

Intro 2 ML

57

Multilayer neural networks



Intro 2 ML

58

Different Non-Linearly Separable Problems

Structure	Types of Decision Regions	Exclusive-OR Problem	Classes with Meshed regions	Most General Region Shapes
Single-Layer 	Half Plane Bounded By Hyperplane			
Two-Layer 	Convex Open Or Closed Regions			
Three-Layer 	Arbitrary (Complexity Limited by No. of Nodes)			

4/10/2024

59


Backpropagation

- Algorithm proposed in 1970.
- Became convincingly popular in 1986 due to a paper by [David Rumelhart](#), [Geoffrey Hinton](#), and [Ronald Williams](#).
- At the core of backpropagation is an expression for the partial derivative of Error function with respect to weights, i.e., $\frac{\partial E}{\partial w}$

Intro 2 ML

60

Weights of neural network



Intro 2 ML

61

Output of j -th node at the l -th layer

- Input to l -th layer is coming from $(l-1)$ -th layer, i.e., $\mathbf{y}^{(l-1)}$.
- Suppose there are K nodes in the $(l-1)$ -th layer.
 - $\mathbf{y}^{(l-1)} = (1, y_1^{(l-1)}, y_2^{(l-1)}, y_3^{(l-1)}, \dots, y_{K-1}^{(l-1)})^T$
- Weight of the connection from k -th node of the $(l-1)$ -th layer to the j -th node of the l -th layer is $w_{jk}^{(l)}$.
 - $\mathbf{w}_j^{(l)} = (w_{j0}^{(l)}, w_{j1}^{(l)}, w_{j2}^{(l)}, \dots, w_{jK-1}^{(l)})^T$, where $w_{j0}^{(l)}$ is the weight to the bias.

Intro 2 ML

62

Output of j -th node at the l -th layer

- Output of the j -th node at the l -th layer is

$$y_j^{(l)} = \sigma \left(\left(\mathbf{w}_j^{(l)} \right)^T \mathbf{y}^{(l-1)} \right)$$

where $l = 1, 2, 3, \dots, L$ and that means the NN has $L-1$ hidden layers.

- Note that at the input layer, i.e., $\mathbf{y}^{(0)} = \mathbf{x}$ and output is $\mathbf{y}^{(L)} = \hat{\mathbf{y}}$.

- Let us decompose $y_j^{(l)} = \sigma \left(\left(\mathbf{w}_j^{(l)} \right)^T \mathbf{y}^{(l-1)} \right)$ into

$$y_j^{(l)} = \sigma \left(z_j^{(l)} \right) \text{ where } z_j^{(l)} = \sum_k w_{jk}^{(l)} y_k^{(l-1)}$$

Intro 2 ML

63

Chain rule to compute

- Considering single output node, rewrite $y_j^{(l)} = \sigma \left(\left(\mathbf{w}_j^{(l)} \right)^T \mathbf{y}^{(l-1)} \right)$ as

$$y_j^{(l)} = \sigma \left(z_j^{(l)} \right) \text{ where } z_j^{(l)} = \sum_k w_{jk}^{(l)} y_k^{(l-1)}$$

- Following the chain rule:

$$\begin{aligned} \hat{y}(\mathbf{x}, \mathbf{w}) = y^{(L)} &= \sigma \left(\sum_k w_{jk}^{(L)} y_k^{(L-1)} \right) \\ &= \sigma \left(\sum_k w_{jk}^{(L)} \sigma \left(z_k^{(L-1)} \right) \right) \\ &= \sigma \left(\sum_k w_{jk}^{(L)} \sigma \left(\sum_m w_{km}^{(L-1)} y_m^{(L-2)} \right) \right) \dots \end{aligned}$$

Intro 2 ML

64

Derivative of function of functions

- Suppose we have $f(x) = \log_e(\sin(x^2))$
- Consider $f(x) = f_1(y)$ where $y = \sin(x^2)$
- then $y = f_2(z)$ where $z = x^2$
- then $z = f_3(x)$
- Thus $f(x) = f_1(y) \Rightarrow f(x) = f_1(f_2(z)) \Rightarrow f(x) = f_1(f_2(f_3(x)))$
- $\frac{df}{dx} = \frac{df_1}{df_2} \frac{df_2}{df_3} \frac{df_3}{dx}$ OR $\frac{df}{dx} = \frac{df}{dy} \frac{dy}{dz} \frac{dz}{dx}$ OR $\frac{df}{dx} = \frac{1}{\sin(x^2)} \cos(x^2) 2x$

Intro 2 ML

65

Backpropagation

- We do not have groundtruth at the output of every layer, except the final layer, change in weight at any layer is related to the change in error ΔE as

$$\Delta E = \frac{\partial E}{\partial w_{jk}^{(l)}} \Delta w_{jk}^{(l)}$$

Recall that

$$\hat{y}(\mathbf{x}, \mathbf{w}) = y^{(L)} = \sigma \left(\sum_k w_{jk}^{(L)} \sigma \left(\sum_m w_{km}^{(L-1)} y_m^{(L-2)} \right) \right) \dots$$

Intro 2 ML

66

Backpropagation (contd.)

- However, to compute total change in error ΔE due to change in weights of k -th node of $(l-1)$ -th layer connected to the j -th node of l -th layer, it is plausible that we should sum over all possible paths from k -th node of the $(l-1)$ -th layer to the final layer, i.e.,

$$\Delta E \approx \sum_{mnp \dots qr} \frac{\partial E}{\partial y_m^{(L)}} \frac{\partial y_m^{(L)}}{\partial y_n^{(L-1)}} \frac{\partial y_n^{(L-1)}}{\partial y_p^{(L-2)}} \dots \frac{\partial y_q^{(l+1)}}{\partial y_r^{(l)}} \frac{\partial y_r^{(l)}}{\partial w_{jk}^{(l)}} \Delta w_{jk}^{(l)}$$

Intro 2 ML

67

Backpropagation (contd.)

- Now combining following two equations:

$$\Delta E = \frac{\partial E}{\partial w_{jk}^{(l)}} \Delta w_{jk}^{(l)} \text{ and}$$

$$\Delta E \approx \sum_{mnp \dots qr} \frac{\partial E}{\partial y_m^{(L)}} \frac{\partial y_m^{(L)}}{\partial y_n^{(L-1)}} \frac{\partial y_n^{(L-1)}}{\partial y_p^{(L-2)}} \dots \frac{\partial y_q^{(l+1)}}{\partial y_r^{(l)}} \frac{\partial y_r^{(l)}}{\partial w_{jk}^{(l)}} \Delta w_{jk}^{(l)}$$

- We obtain

$$\frac{\partial E}{\partial w_{jk}^{(l)}} = \sum_{mnp \dots qr} \frac{\partial E}{\partial y_m^{(L)}} \frac{\partial y_m^{(L)}}{\partial y_n^{(L-1)}} \frac{\partial y_n^{(L-1)}}{\partial y_p^{(L-2)}} \dots \frac{\partial y_q^{(l+1)}}{\partial y_r^{(l)}} \frac{\partial y_r^{(l)}}{\partial w_{jk}^{(l)}}$$

Intro 2 ML

68

Updating weight

- Error function: $E(\mathbf{w}) = \frac{1}{2n} \sum_x ||\hat{y}(\mathbf{x}, \mathbf{w}) - y(\mathbf{x})||^2$

where $\hat{y}(\mathbf{x}, \mathbf{w}) = \sigma \left(\sum_k w_{jk}^{(l)} \sigma \left(\sum_m w_{km}^{(l-1)} y_m^{(l-2)} \right) \right) \dots$

- Earlier we had (for single layer): $w_k^{(t+1)} = w_k^{(t)} - \eta \frac{\partial E}{\partial w_k}$
- Now updating weight from k -th node of $(l-1)$ -th layer to j -th node of l -th layer as

$$w_{jk}^{(l)(t+1)} = w_{jk}^{(l)(t)} - \eta \frac{\partial E}{\partial w_{jk}^{(l)(t)}}$$

Intro 2 ML

69

Standard activation functions

- Sigmoid functions ('S'-shaped curves)
 - The hyperbolic tangent (symmetrical)
 - Both functions have a simple differential
 - Only the shape is important (**Neuron**)


$$\sigma(z) = \frac{1}{1 + e^{-az}}$$

4/10/2024

70

Sigmoid function

- Sigmoid function
 - may be expressed as $\sigma(z) = \frac{1}{1 + e^{-az}}$
 - is one of the most popular activation function.
 - squashes the input value between 0 and 1.
 - is smooth and differentiable.
 - maximum slope is at $z = 0$



Intro 2 ML

71

Derivative of sigmoid function


- We have $y = \sigma(z) = \frac{1}{1 + e^{-az}}$
- Derivative of $\sigma(z)$ at $z = 0$ may be written as

$$\frac{d\sigma}{dz} \Big|_{z=0} = \frac{e^{-az}}{(1 + e^{-az})^2} \Big|_{z=0} = \frac{1}{(1 + 1)^2} = 0.25$$
- This is the maximum value of gradient for any z .

Intro 2 ML

72

Multilayer neural network: Example



73

Vanishing gradient problem

- Number of layers are usually approximates the degree of polynomial function it can realize.
- However, more layers means more neurons and consequently more time to train the network.
- Second, since the derivative of the activation function (resulting in output at each layer) ≤ 0.25 ,

$$\frac{\partial E}{\partial w_{jk}^{(l)}} = \sum_{mnp \dots qr} \frac{\partial E}{\partial y_m^{(L)}} \frac{\partial y_m^{(L)}}{\partial y_n^{(L-1)}} \frac{\partial y_n^{(L-1)}}{\partial y_p^{(L-2)}} \dots \frac{\partial y_q^{(l+1)}}{\partial y_r^{(l)}} \frac{\partial y_r^{(l)}}{\partial w_{jk}^{(l)}}$$

- May tend to zero. This is known as **vanishing gradient problem**.
- This problem is more evident as we deeper layers from output input.

Intro 2 ML

74

Vanishing gradient problem

- Recall the weight updating rule

$$w_{jk}^{(l)(t+1)} = w_{jk}^{(l)(t)} - \eta \frac{\partial E}{\partial w_{jk}^{(l)(t)}}$$


- If $\frac{\partial E}{\partial w_{jk}^{(l)(t)}} \rightarrow 0$, we have $w_{jk}^{(l)(t+1)} \approx w_{jk}^{(l)(t)}$

- The first layers are supposed to carry most of the information, but we see it gets trained the least.
- Hence, the problem of vanishing gradient eventually leads to the death of the network.

Intro 2 ML

75

Vanishing gradient problem



Intro 2 ML

76

Exploding gradient problem


- Suppose vanishing gradient problem does not occur.
- Then $\frac{\partial E}{\partial w_{jk}^{(l)}} = \sum_{mn} \dots q_r \frac{\partial E}{\partial y_m^{(L)}} \frac{\partial y_m^{(L)}}{\partial y_n^{(L-1)}} \frac{\partial y_n^{(L-1)}}{\partial y_p^{(L-2)}} \dots \frac{\partial y_q^{(L+1)}}{\partial y_r^{(l)}} \frac{\partial y_r^{(l)}}{\partial w_{jk}^{(l)}}$ implies that
- $\frac{\partial E}{\partial w_{jk}^{(l)}}$ is a sum of gradient magnitude along $m \times n \times p \times \dots \times q \times r$ number of paths, where each gradient is greater than 0.
- Thus this sum could be significantly high resulting in **exploding gradient problem**.

Activation function (non-linear)

- Rectified Linear Unit (ReLU): $y = \max(0, x)$

- Softplus function: $y = \log(1 + e^x)$

- Leaky ReLU:



Some hyperparameters

- **Epoch:** Suppose there are n samples in the training set. Passing (or using) all n samples to train the network is known as one epoch.
 - To train the network we need pass the training samples over and over again.
 - As the number of epoch increases network upgrades **from underfitting to optimum to overfitting**.
- **Batch:** If n is large, training set is divided into small batches or groups or sets of training data. **Batch size** is the number of training samples, say m , in each batch.
- **Iteration:** The number of batches that are passed through the network to complete one epoch, i.e., n/m .

Batch normalization

- As the training progresses the network encounters (or being feed into) newer data
 - The statistical distribution of the input to layer(s) keeps changing.
 - The distribution of the output of each layer in different batches are different.
 - This reduces training efficiency.
- The input samples (in every batch) are normalized before feeding it into the next layer of the network.
 - The mean and variance of all such batches, instead of the entire data, are computed.
 - This is known as **batch normalization**.

Dropout

- This is used to overcome the overfitting problem.
- Often certain nodes in the network are randomly switched off, from some or all the layers of a neural network.
 - Hence, in every iteration, we get a new network.
 - The resulting network (obtained at the end of training) is a combination of all of them.
 - This is an way of implementing the *regularization*.

Thank you!
Any question?